



Data Wrangling in R - Data Wrangling - 3

One should look for what is and not what he thinks should be. (Albert Einstein)

Warm-up: review activity

- Match each term to the correct definition below, and share responses in the chat (ex: “1,E” to match the first function to the definition labelled “E”)

1. `filter`
2. `arrange`
3. `select`
4. `mutate`
5. `summarize`
6. `group_by`

- A. Pick variables by their names
- B. Collapse many values down to a single summary
- C. Reorder the rows
- D. Allows the above functions to operate on a dataset group by group
- E. Pick observations by their value
- F. Create new variables with functions of existing variables

Warm-up: Review activity answers

- The answers are: **1,E; 2,C; 3,A; 4,F; 5,B; 6,D**

Function	Use Case
<code>filter</code>	Pick observations by their value
<code>arrange</code>	Reorder the rows
<code>select</code>	Pick variables by their names
<code>mutate</code>	Create new variables with functions of existing variables
<code>summarize</code>	Collapse many values down to a single summary
<code>group_by</code>	Allows the above functions to operate on a dataset group by group

- Any questions?

Module completion checklist

Objective	Complete
Apply the filter function to subset data	
Rank data using the arrange function	

Getting started with dplyr

- Now that we have dplyr installed and the nycflights13 dataset loaded, we can start transforming our dataset
- Our goal is to get exposure to the core dplyr syntax and practice using some of the major function verbs
- For now, we will focus on just two:
 - `filter()`, to subset the observations in the data
 - `arrange()`, to reorder the observations in the data

Filter

- `filter` allows you to subset observations based on their values
- Basic use cases for `filter` function include:
 - Subsetting the data to include flights from January 2013
 - Subsetting the data that contain missing values

```
# Check for detailed documentation
?dplyr::filter

# Use cases for `filter` function

filter(df,                #<- dataframe
       filter_cond1,     #<- subsetting rule(s)
       ...)              #<- other arguments
```

filter {dplyr}

R Documentation

Subset rows using column values

Description

The `filter()` function is used to subset a data frame, retaining all rows that satisfy your conditions. To be retained, the row must produce a value of `TRUE` for all conditions. Note that when a condition evaluates to `NA` the row will be dropped, unlike base subsetting with `[]`.

Usage

```
filter(.data, ..., .preserve = FALSE)
```

Arguments

<code>.data</code>	A data frame, data frame extension (e.g. a tibble), or a lazy data frame (e.g. from <code>dbplyr</code> or <code>dtplyr</code>). See <i>Methods</i> , below, for more details.
--------------------	---

Filter

- Let's say you would like to see all flights from January 2013
- Using `filter`, pass the original dataframe followed by filtering criteria

```
# Load the flights dataset into the environment.
data(flights)

# Filter `flights` data frame to display all records from January (month == 1) of 2013 (year == 2013).
filter(flights,      #<- set data
       month == 1,    #<- filter by month
       year == 2013)  #<- filter by year
```

```
# A tibble: 27,004 x 19
   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
   <int> <int> <int>   <int>         <int>         <dbl>     <int>         <int>
1  2013     1     1     517           515           2       830           819
2  2013     1     1     533           529           4       850           830
3  2013     1     1     542           540           2       923           850
4  2013     1     1     544           545          -1      1004          1022
5  2013     1     1     554           600          -6       812           837
6  2013     1     1     554           558          -4       740           728
7  2013     1     1     555           600          -5       913           854
8  2013     1     1     557           600          -3       709           723
9  2013     1     1     557           600          -3       838           846
10 2013     1     1     558           600          -2       753           745
# ... with 26,994 more rows, and 11 more variables: arr_delay <dbl>,
#   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

Filter

- If you want to build on top of the filtered dataset, you will need to save your new subset to a new variable and perform further operations on this new subset

```
# You will have to make sure to save the subset. To do this, use `=`.  
filter_flights = filter(flights, month == 1, year == 2013)  
  
# View your output.  
filter_flights
```

```
# A tibble: 27,004 x 19  
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time  
  <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>  
1  2013     1     1     517           515         2      830           819  
2  2013     1     1     533           529         4      850           830  
3  2013     1     1     542           540         2      923           850  
4  2013     1     1     544           545        -1     1004          1022  
5  2013     1     1     554           600        -6      812           837  
6  2013     1     1     554           558        -4      740           728  
7  2013     1     1     555           600        -5      913           854  
8  2013     1     1     557           600        -3      709           723  
9  2013     1     1     557           600        -3      838           846  
10 2013     1     1     558           600        -2      753           745  
# ... with 26,994 more rows, and 11 more variables: arr_delay <dbl>,  
#   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,  
#   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```


Filter options

- You can use the standard filtering operations when working with integer data types:


Operation	Use Case	Example
>	Greater than	6 > 4
>=	Greater than or equal to	4 >= 4
<	Less than	4 < 6
<=	Less than or equal to	4 <= 4
!=	Not equal to	4 != 6
==	Equal to	4 == 4

Filter options (cont'd)

- And more general operators:

Operation	Use Case	Example
	either can be true to satisfy	<code>x == 4 x == 12, x==2 x==13</code>
&	and, both need to be true	<code>x == 4 & y == 2</code>
!	Not true, inverse selection	<code>x != 4</code>
<code>%in%</code>	value in the following list of values	<code>x %in% c(4,16,32)</code>

Filter - logical operators



Filter - examples of logical operators

- What if we want to see all flights from January **and** on the 25th?

```
# Filter with just `&`.  
filter(flights, month == 1 & day == 25)
```

```
# A tibble: 922 x 19  
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time  
  <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>  
1  2013     1    25      15          1815          360     208          1958  
2  2013     1    25      17          2249           88     119          2357  
3  2013     1    25      26          1850          336     225          2055  
4  2013     1    25     123          2000          323     229          2101  
5  2013     1    25     123          2029          294     215          2140  
6  2013     1    25     456           500           -4     632           648  
7  2013     1    25     519           525           -6     804           820  
8  2013     1    25     527           530           -3     820           829  
9  2013     1    25     535           540           -5     826           850  
10 2013     1    25     539           540           -1    1006          1017  
# ... with 912 more rows, and 11 more variables: arr_delay <dbl>, carrier <chr>,  
# flight <int>, tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,  
# distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

- Note:** After running each example, we will record the number of rows. This will help illustrate each operator and how different a simple change of one boolean operator can have on the dataset. Total number of rows should be 922

Filter - examples of logical operators (cont'd)

- What if we want to see all flights, but **exclude** those from January and those on the 25th?

```
# Filter with `!`.  
filter(flights, month != 1 & day != 25)
```

```
# A tibble: 299,597 x 19  
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time  
  <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>  
1  2013     10      1     447             500         -13     614             648  
2  2013     10      1     522             517          5     735             757  
3  2013     10      1     536             545         -9     809             855  
4  2013     10      1     539             545         -6     801             827  
5  2013     10      1     539             545         -6     917             933  
6  2013     10      1     544             550         -6     912             932  
7  2013     10      1     549             600        -11     653             716  
8  2013     10      1     550             600        -10     648             700  
9  2013     10      1     550             600        -10     649             659  
10 2013     10      1     551             600         -9     727             730  
# ... with 299,587 more rows, and 11 more variables: arr_delay <dbl>,  
#   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,  
#   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

- Here we are looking for all flights that are not in January and not on the 25th; total number of rows should be 299,597

Filter - examples of logical operators (cont'd)

```
# Filter with `%in%`.
filter(flights, month %in% c(1, 2) & day == 25)
```

```
# A tibble: 1,883 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
1  2013     1    25      15          1815         360     208          1958
2  2013     1    25      17          2249          88     119          2357
3  2013     1    25      26          1850         336     225          2055
4  2013     1    25     123          2000         323     229          2101
5  2013     1    25     123          2029         294     215          2140
6  2013     1    25     456           500          -4     632           648
7  2013     1    25     519           525          -6     804           820
8  2013     1    25     527           530          -3     820           829
9  2013     1    25     535           540          -5     826           850
10 2013     1    25     539           540          -1    1006          1017
# ... with 1,873 more rows, and 11 more variables: arr_delay <dbl>,
#   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

- This is a combination of `&` and `%in%` subsetting all flights from January and February that are on the 25th; number of rows should be 1,883

Using filter with NA values

- `filter` only includes rows where the condition is TRUE; it **excludes** both FALSE and NA values
- If you want to preserve missing values, ask for them explicitly

```
# Create a data frame with 2 columns.  
NA_df = data.frame(x = c(1, NA, 2), #<- column x with 3 entries with 1 NA  
                   y = c(1, 2, 3))   #<- column y with 3 entries  
  
# Filter without specifying anything regarding NAs.  
filter(NA_df, x >= 1)
```

```
  x y  
1 1 1  
2 2 3
```

```
# Filter with specifying to keep rows if there is an NA.  
filter(NA_df, is.na(x) | x >= 1)
```

```
  x y  
1 1 1  
2 NA 2  
3 2 3
```

Module completion checklist

Objective	Complete
Apply the filter function to subset data	✓
Rank data using the arrange function	

Arrange

- `arrange` is used to reorder the observations within the specified column(s)
- It is the equivalent of `sort` in SAS or `order by` in SQL

```
# Check for detailed documentation
?dplyr::arrange

# Use cases for `arrange` function
arrange(df,                #<- data frame
         arrange_cond1,    #<- column by which
                           #   to arrange
         ...)              #<- other args.
```

`arrange {dplyr}`

R Documentation

Arrange rows by column values

Description

`arrange()` order the rows of a data frame rows by the values of selected columns.

Unlike other dplyr verbs, `arrange()` largely ignores grouping; you need to explicitly mention grouping variables (or use `by_group = TRUE`) in order to group by them, and functions of variables are evaluated once per data frame, not once per group.

Usage

```
arrange(.data, ..., .by_group = FALSE)
```

```
## S3 method for class 'data.frame'
arrange(.data, ..., .by_group = FALSE)
```

Arguments

<code>.data</code>	A data frame, data frame extension (e.g. a tibble), or a lazy data frame (e.g. from <code>dbplyr</code> or <code>dtplyr</code>). See <i>Methods</i> , below, for more details.
<code>...</code>	<data-masking> Variables, or functions or variables. Use desc() to sort a variable in descending order.
<code>.by_group</code>	If <code>TRUE</code> , will sort first by grouping variable. Applies to grouped data frames only.

Arrange example

- When using multiple columns with `arrange`, the additional columns will be used to break ties in the values of preceding columns

```
# Arrange data by year, then month, and then day.
arrange(flights, #<- data frame we want to arrange
  year,         #<- 1st: arrange by year
  month,        #<- 2nd: arrange by month
  day)          #<- 3rd: arrange by day
```

```
# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
1  2013     1     1     517             515         2       830           819
2  2013     1     1     533             529         4       850           830
3  2013     1     1     542             540         2       923           850
4  2013     1     1     544             545        -1      1004          1022
5  2013     1     1     554             600        -6       812           837
6  2013     1     1     554             558        -4       740           728
7  2013     1     1     555             600        -5       913           854
8  2013     1     1     557             600        -3       709           723
9  2013     1     1     557             600        -3       838           846
10 2013     1     1     558             600        -2       753           745
# ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
#   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

Arrange options

- `arrange` by default sorts everything in ascending order; to arrange in descending, use **`desc`**
- You can now see that the month at the top of the dataset is **December** (i.e., 12th month)

```
# Arrange data by year, descending month and then day.
arrange(flights,      #<- data frame we want to arrange
        year,        #<- 1st: arrange by year
        desc(month), #<- 2nd: arrange by month in descending order
        day)         #<- 3rd: arrange by day
```

```
# A tibble: 336,776 x 19
   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
1  2013     12     1      13           2359           14     446           445
2  2013     12     1      17           2359           18     443           437
3  2013     12     1     453           500           -7     636           651
4  2013     12     1     520           515            5     749           808
5  2013     12     1     536           540           -4     845           850
6  2013     12     1     540           550          -10    1005          1027
7  2013     12     1     541           545           -4     734           755
8  2013     12     1     546           545            1     826           835
9  2013     12     1     549           600          -11     648           659
10 2013     12     1     550           600          -10     825           854
# ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
#   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

Arrange with NA values

- Missing values are **always** sorted at the end

```
# Arrange data with missing values.  
arrange(NA_df, x)
```

	x	y
1	1	1
2	2	3
3	NA	2

```
# Even when we use `desc` the `NA` is taken to the last row.  
arrange(NA_df, desc(x))
```

	x	y
1	2	3
2	1	1
3	NA	2

Knowledge check



Module completion checklist

Objective	Complete
Apply the filter function to subset data	✓
Rank data using the arrange function	✓

Congratulations on completing this module!

You are now ready to try Tasks 6-12 in the Exercise for this topic

